

Session 9, 10, 11 - Searching and Sorting Algorithms



Objective of Searching:

The objective of searching is to find relevant information, answers, or solutions to a specific query or problem. Searching is a fundamental human activity that has been significantly facilitated by the rise of the internet and search engines. The objectives of searching can vary depending on the context, but some common objectives include:

- **Information retrieval:** Searching aims to retrieve accurate and relevant information from various sources such as websites, databases, books, articles, or any other digital or physical resources.
- **Problem-solving:** People often search for solutions to specific problems they encounter, such as troubleshooting technical issues, finding ways to fix or repair something, or seeking advice on a particular challenge.
- **Knowledge acquisition:** Searching allows individuals to acquire new knowledge on a wide range of topics, from academic subjects to hobbies, history, current events, and more.
- **Decision-making:** When faced with choices or decisions, searching can provide valuable insights and data to help individuals make informed and well-reasoned judgments.
- **Validation and verification:** People may search to validate or verify information they have come across to ensure its accuracy and credibility.
- **Exploration and discovery:** Searching can be an exploratory activity to discover new ideas, perspectives, products, or services.
- **Communication and social interaction:** Searching can also be used to find and connect with like-minded individuals or communities who share similar interests or opinions.
- **Entertainment and leisure:** Searching for entertainment purposes can involve looking up movies, games, music, or other forms of media to enjoy in one's leisure time.
- **Academic and professional research:** Students, academics, and professionals search to gather information for research, papers, reports, or projects.

- **Personal growth and self-improvement:** Individuals might search for self-help resources, educational content, or motivational materials to aid in personal development and growth.


Overall, the objectives of searching revolve around obtaining relevant and reliable information to meet specific needs, solve problems, and enrich our understanding of the world around us. Effective searching involves using appropriate search strategies and tools to achieve the desired outcomes efficiently.

Sequential Search

Sequential search, also known as linear search, is a simple searching algorithm used to find a specific target value within a list or array of elements. It sequentially checks each element in the list from the beginning until it finds the target element or reaches the end of the list.

Here's how the sequential search algorithm works:

- Start at the first element of the list.
- Compare the target value with the current element.
- If the current element matches the target value, the search is successful, and the algorithm returns the index or position of the element in the list.
- If the current element does not match the target value, move to the next element in the list and repeat steps 2 and 3.
- Continue this process until the target value is found, or all elements have been checked.



If the target value is not present in the list, the algorithm will reach the end of the list without finding a match and will return a "not found" indication.

Analysis of Sequential Search algorithm

To analyse the sequential search algorithm, we'll examine its time complexity and space complexity.

Time Complexity:

- The time complexity of the sequential search algorithm is $O(n)$, where n is the number of elements in the list. In the worst-case scenario, the algorithm will have to compare the target element with each element in the list until it finds a match or reaches the end of the list. This means that the time taken to execute the algorithm grows linearly with the size of the input.

In the best-case scenario, the target element is found at the very first position, resulting in only one comparison. However, best-case analysis is often not as informative as worst-case analysis, especially when we need to consider the algorithm's performance in general cases.

It's essential to note that sequential search does not take advantage of the list's ordering; hence, it needs to check every element one by one. Therefore, for very large lists, the time complexity of $O(n)$ can become a performance bottleneck.

Space Complexity:

- The space complexity of the sequential search algorithm is $O(1)$ since it uses a constant amount of additional memory to store variables like the loop index and temporary variables. The space required does not depend on the size of the input list, making it an in-place algorithm.

Advantages of Sequential Search:

- **Simplicity:** The algorithm is straightforward to understand and implement, making it a good choice for small lists or when the list is unsorted.
- **In-place:** The algorithm doesn't require additional memory proportional to the size of the input list, making it memory-efficient.

Disadvantages of Sequential Search:

- **Inefficiency:** The time complexity of $O(n)$ can be inefficient for large lists, especially when faster search algorithms like binary search are available for sorted lists.
- **Lack of optimization:** The algorithm doesn't take advantage of any special characteristics of the list, such as being sorted, which could lead to suboptimal performance.

Use Cases:

Sequential search is best suited for small lists or unsorted lists where other more efficient search algorithms might not offer significant advantages. It is also useful when a simple search algorithm is needed in situations where the data set is relatively small or dynamic.

However, when dealing with large sorted lists or databases, more advanced search algorithms like binary search or hash-based techniques are preferred as they offer logarithmic or constant time complexity, respectively, which can significantly improve search efficiency.


Binary Search

Binary search is an efficient search algorithm used to find a specific target value within a sorted list or array. It works by repeatedly dividing the search space in half and eliminating the half where the target value cannot be present. Binary search is a logarithmic algorithm, meaning its time complexity grows logarithmically with the size of the input data.

Here's how the binary search algorithm works:

- Start with the entire sorted list (or array).
- Calculate the middle index of the current search space.
- Compare the middle element with the target value:
 - If the middle element matches the target, the search is successful, and the algorithm returns the index or position of the element.
 - If the middle element is greater than the target, discard the right half of the search space and continue the search in the left half.
 - If the middle element is less than the target, discard the left half of the search space and continue the search in the right half.
- Repeat steps 2 and 3 until the target value is found or the search space is empty.

Binary search is particularly efficient because it halves the search space at each step, which reduces the number of elements to check significantly with each iteration.



Analysis of Binary Search

Let's analyze the binary search algorithm in terms of its time complexity, space complexity, and other considerations.

Time Complexity:

- The time complexity of binary search is $O(\log n)$, where n is the number of elements in the sorted list or array. With each comparison, the search space is halved, leading to a significant reduction

in the number of elements to check at each iteration. This logarithmic behavior makes binary search very efficient, especially for large data sets.

In each iteration, the size of the search space is reduced to approximately half of its previous size. As a result, binary search exhibits excellent performance even for very large lists, as the number of elements to be checked decreases rapidly, reducing the search time logarithmically.

Space Complexity:

- The space complexity of binary search is $O(1)$ since the algorithm uses only a constant amount of additional memory to store variables like the left and right pointers and the middle index. The space required does not depend on the size of the input list, making it an in-place algorithm.

Sorted Data Requirement:

- One of the key requirements for binary search is that the input list must be sorted in ascending or descending order. If the list is not sorted, binary search cannot be applied. Sorting the list itself requires at least $O(n \log n)$ time complexity in most cases, which may impact the overall efficiency of the search. Therefore, binary search is best suited for scenarios where the data is already sorted or can be efficiently sorted once and then searched multiple times.

Strengths and Limitations:

- Binary search is highly efficient and performs well on sorted lists, making it a preferred choice for searching operations in such scenarios. However, its main limitation is that it requires the list to be sorted, which can be an overhead if the data is frequently changing or if it is not feasible to maintain a sorted list.

Additionally, binary search may not be the best choice when dealing with unsorted or dynamically changing data, as its sorting requirement and its logarithmic behavior may not be fully utilized in such cases.

Search vs. Preprocessing:

- In situations where the data is frequently changing and needs to be searched multiple times, binary search may be efficient if the preprocessing cost of sorting is not significant compared to the overall search operations. For example, if you have a large sorted dataset and need to perform numerous searches, the initial sorting cost might be justified by the subsequent fast search times.

In summary, binary search is an efficient search algorithm with a time complexity of $O(\log n)$. It is best suited for scenarios where the data is already sorted or can be efficiently sorted once and then searched multiple times. For unsorted or frequently changing data, other search algorithms like sequential search might be more appropriate.

Introduction to sorting

Sorting is the process of arranging elements in a specific order, typically in ascending or descending order, based on certain criteria or key values. Sorting is a fundamental and commonly used operation in computer science and programming. It plays a crucial role in various applications and algorithms, allowing data to be organized systematically and efficiently for easier retrieval, analysis, and processing.

The primary objectives of sorting are:

- **Order:** Sorting arranges elements in a well-defined order, making it easier for humans to understand and for computers to process data more efficiently.
- **Searching:** Sorted data allows for faster search operations using techniques like binary search, which take advantage of the ordered arrangement to locate elements more quickly.
- **Duplicate Removal:** Sorting simplifies the process of identifying and removing duplicate elements from a dataset, as duplicates are likely to be adjacent to each other in the sorted order.
- **Data Analysis:** In various data analysis tasks, sorting helps identify patterns, trends, and anomalies in the data, as it brings similar elements together.
- **Preprocessing:** Sorting often serves as a preliminary step in many algorithms and applications, preparing the data for subsequent operations.

There are numerous sorting algorithms, each with its unique approach, advantages, and limitations. Some commonly used sorting algorithms include:

- **Bubble Sort:** A simple comparison-based sorting algorithm that repeatedly swaps adjacent elements if they are in the wrong order until the entire list is sorted.
- **Selection Sort:** Another simple comparison-based sorting algorithm that repeatedly selects the minimum (or maximum) element and places it in the correct position.
- **Insertion Sort:** A comparison-based sorting algorithm that builds the final sorted list one item at a time by inserting each element into its proper position.
- **Merge Sort:** A divide-and-conquer sorting algorithm that recursively divides the list into smaller sublists, sorts them individually, and then merges them to produce a sorted output.
- **Quick Sort:** A highly efficient divide-and-conquer sorting algorithm that selects a pivot element, partitions the data into two sublists, and recursively sorts them.
- **Heap Sort:** A comparison-based sorting algorithm that uses a binary heap data structure to sort elements efficiently.

Each sorting algorithm has its own strengths and weaknesses, and the choice of the appropriate algorithm depends on the specific use case, the size of the data, and the desired performance characteristics.

Sorting is a fundamental concept in computer science, and understanding different sorting algorithms can greatly enhance a programmer's problem-solving skills and algorithmic understanding.

Selection Sort


Selection Sort is a simple comparison-based sorting algorithm that works by repeatedly selecting the smallest (or largest) element from the unsorted part of the list and swapping it with the first unsorted element. The algorithm divides the list into two parts: the left part, which is sorted, and the right part, which is unsorted. With each iteration, it expands the sorted part by one element until the entire list becomes sorted.

Here's how the Selection Sort algorithm works:

- Start with the first element as the "minimum" (or "maximum") element in the unsorted part of the list.
- Compare the "minimum" element with each element in the unsorted part of the list.

- If an element is found that is smaller (or larger) than the current "minimum" element, update the "minimum" to be the index of that element.
- After iterating through the unsorted part of the list, swap the "minimum" element with the first element in the unsorted part, effectively adding it to the sorted part.
- Repeat steps 1 to 4 until the entire list becomes sorted.

The selection sort algorithm does not make use of any auxiliary data structures and sorts the elements in-place, meaning it only requires a constant amount of additional memory.



Insertion Sort


Insertion Sort is a simple comparison-based sorting algorithm that builds the final sorted list one element at a time. It works by repeatedly taking an element from the unsorted part of the list and inserting it into its correct position in the sorted part of the list. The algorithm divides the list into two parts: the left part, which is sorted, and the right part, which is unsorted. With each iteration, it expands the sorted part by one element until the entire list becomes sorted.

Here's how the Insertion Sort algorithm works:

1. Start with the second element (index 1) as the current element to be inserted into the sorted part of the list.
2. Compare the current element with the elements in the sorted part of the list, moving from right to left, until the correct position for the current element is found.
3. Shift the larger elements to the right to make space for the current element to be inserted in its correct position.
4. Repeat steps 2 and 3 for each subsequent unsorted element until the entire list becomes sorted.

Unlike some other sorting algorithms that use swapping to rearrange elements, Insertion Sort shifts elements to their correct positions, which makes it efficient for partially sorted or small lists.

Insertion Sort Execution Example

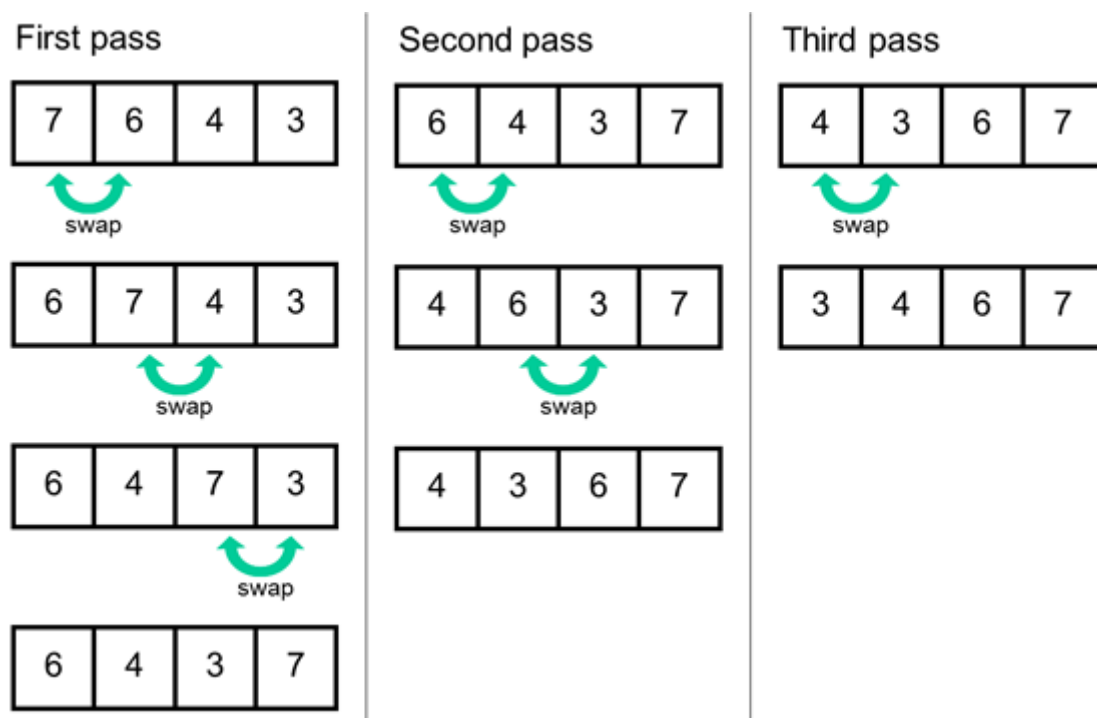


Bubble Sort

Bubble Sort is a simple sorting algorithm that repeatedly steps through the list to be sorted, compares adjacent elements, and swaps them if they are in the wrong order. It gets its name from the way smaller elements "bubble" to the top of the list during each pass.

Here's an explanation of how Bubble Sort works:

1. Start with the first element (index 0) and compare it with the next element (index 1). If the first element is greater than the second element, swap them. Otherwise, leave them in their original order.
2. Move on to the next pair of adjacent elements (indexes 1 and 2) and perform the same comparison and swapping process if necessary. Continue this process, comparing and swapping adjacent elements as you move through the list.
3. After one pass through the list, the largest element will have "bubbled up" to the end of the list (the last index). This means that the largest element is now in its correct sorted position.
4. Repeat steps 1 to 3 for the remaining unsorted part of the list (all elements except the last one, which is already in its correct position). With each pass, one more element will be correctly positioned at the end of the list.
5. Continue this process until the entire list becomes sorted. In each pass, the largest remaining element is correctly placed at its appropriate position.
6. The algorithm includes an optimization to stop early if no swaps were made during a pass. If no swaps occurred, it means the list is already sorted, and the algorithm can terminate early.




Bubble Sort is straightforward to understand and implement, but it has a time complexity of $O(n^2)$ in the worst and average cases, where n is the number of elements in the list. This quadratic time complexity makes it inefficient for large lists, and other more advanced sorting algorithms like Merge Sort or Quick Sort are generally preferred for better performance. Bubble Sort is mainly used for educational purposes or when dealing with very small datasets where simplicity and ease of implementation outweigh the efficiency considerations.

HeapSort

Heap Sort is a comparison-based sorting algorithm that uses the binary heap data structure to sort elements efficiently. It involves two main steps: building a max-heap from the input array and repeatedly extracting the maximum element from the heap and placing it at the end of the array. By doing this, the elements end up in ascending order.

Here's an explanation of how Heap Sort works:

1. **Build Max-Heap:** The first step is to build a max-heap from the input array. A max-heap is a binary tree where the value of each node is greater than or equal to the values of its children. The root of the heap will always hold the maximum element of the array.
2. **Extract Max and Re-heapify:** After building the max-heap, the maximum element (at the root) is swapped with the last element of the heap (the last element of the unsorted part of the array). Then, the heap size is reduced by one (effectively removing the last element from the heap) and the max-heap property is restored by "sifting down" the root element to its proper position.
3. **Repeat:** Steps 2 are repeated until the heap contains only one element. This process extracts the maximum element repeatedly and places it at the end of the array.
4. **Reverse the Array:** After the heap has been fully extracted, the elements are sorted in descending order. To obtain the final sorted array in ascending order, the array is reversed.



Heap Sort has a time complexity of $O(n \log n)$ in all cases, where n is the number of elements in the array. The build max-heap operation takes $O(n)$ time, and the extraction and re-heapification process take $O(\log n)$ time, both of which are performed n times in the worst case. The time complexity is guaranteed to be $O(n \log n)$ regardless of the initial order of the elements, making Heap Sort a reliable and efficient sorting algorithm.

Heap Sort is an in-place sorting algorithm, meaning it uses only a constant amount of additional memory for its operations, making it memory-efficient. However, it is not stable, meaning that the relative order of equal elements may change during the sorting process.

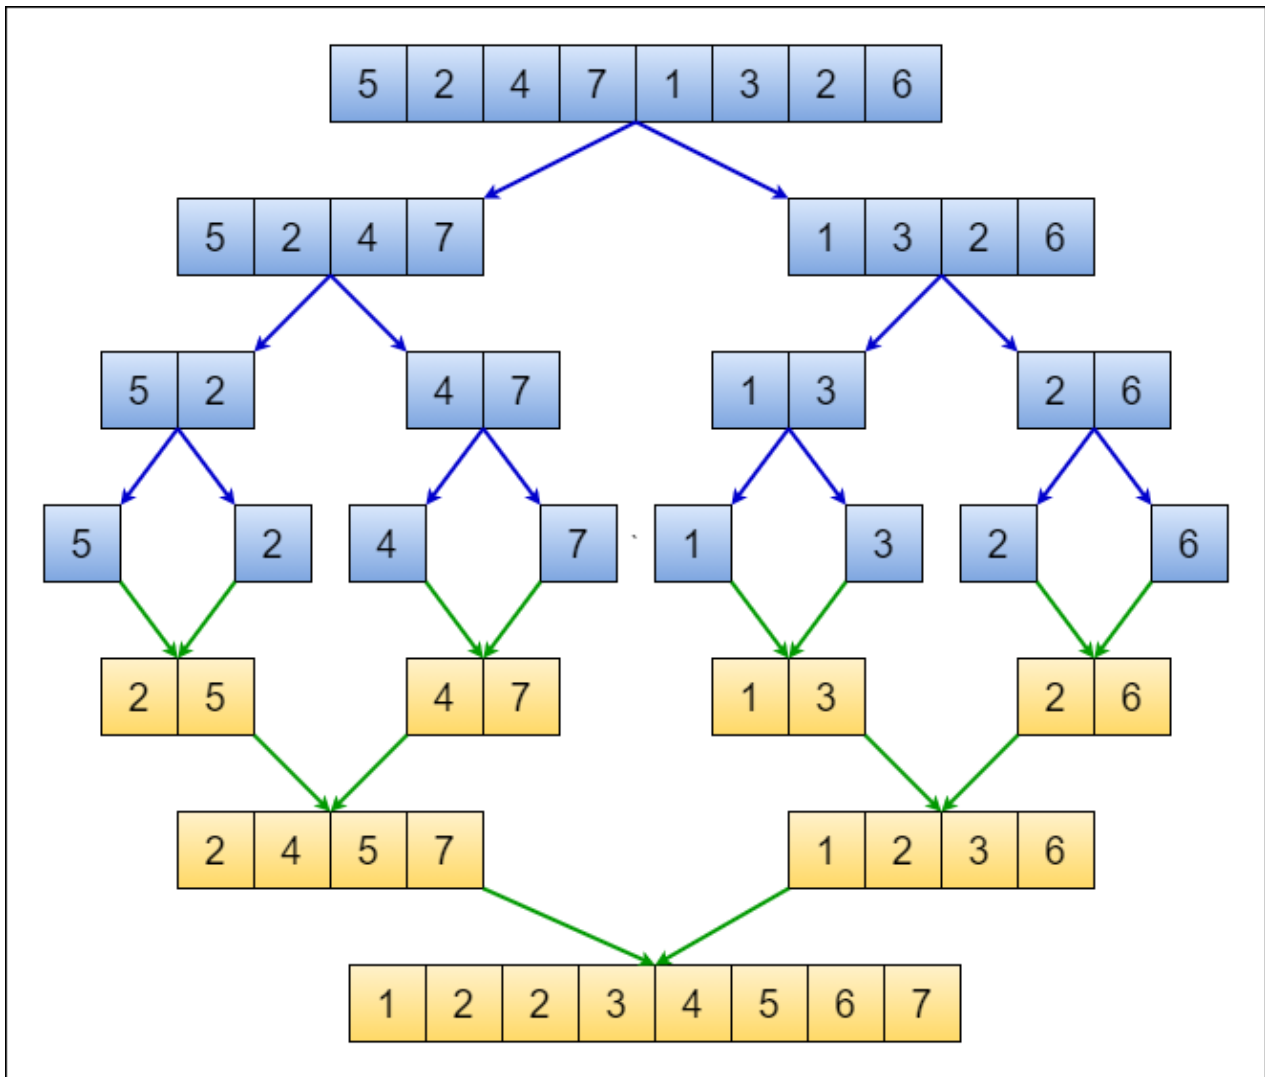
Overall, Heap Sort is a versatile and efficient sorting algorithm suitable for various applications, especially when stable sorting is not a requirement, and the desire for an in-place algorithm is essential.

Merge Sort

Merge Sort is a popular comparison-based sorting algorithm that follows the divide-and-conquer approach. It divides the unsorted list into smaller sublists, recursively sorts each sublist, and then merges them back together to produce a sorted output. Merge Sort guarantees a stable sort, meaning the relative order of equal elements remains unchanged.

Here's an explanation of how Merge Sort works:

1. **Divide:** The unsorted list is recursively divided into two halves until each sublist contains only one element. This step is performed recursively until the base case is reached, i.e., when a sublist has a single element (a single element is always considered sorted).
2. **Conquer (Sort):** After the divide step, the recursion starts to "conquer" the sublists by merging them back together in sorted order. This is achieved by repeatedly merging two sublists while sorting them. During the merging process, elements from both sublists are compared, and the smaller (or larger, depending on the desired order) element is placed in the final merged list. This process continues until all the elements are merged back into a single sorted list.
3. **Combine (Merge):** The merging process involves combining the sorted sublists to form a larger sorted list. This step is essential for achieving the final sorted output.



Merge Sort has a time complexity of $O(n \log n)$ in all cases, where n is the number of elements in the list. The divide and conquer steps both have logarithmic time complexity, and the merging step has linear time complexity, making Merge Sort an efficient sorting algorithm for large datasets.

Merge Sort is also an adaptive sorting algorithm, meaning that its performance can be optimized if the list is partially sorted. Additionally, it is a stable sort, preserving the relative order of equal elements.

While Merge Sort is efficient and reliable, it does require additional memory for its operations. The merging step involves creating additional temporary arrays, which contributes to the space complexity of $O(n)$ in the worst case. As a result, Merge Sort might not be the best choice for sorting very large datasets with limited memory.

Overall, Merge Sort is a powerful and widely used sorting algorithm due to its efficiency, stability, and ability to handle various types of data.

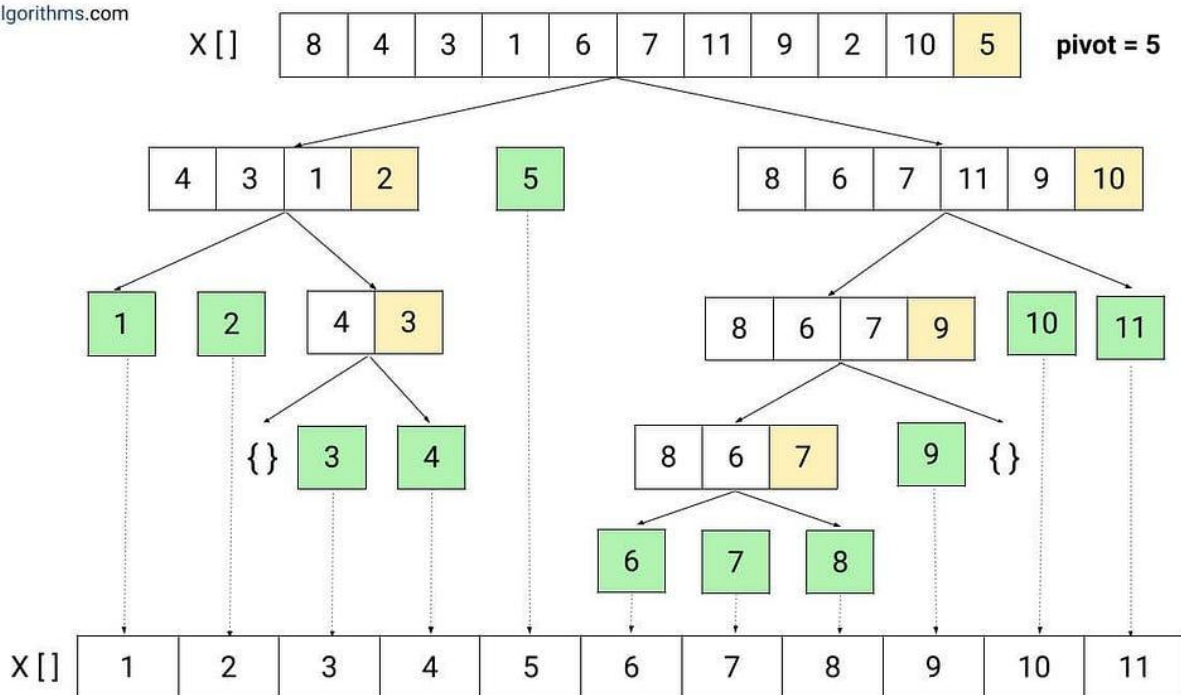
Quick Sort

Quick Sort is a widely used comparison-based sorting algorithm that follows the divide-and-conquer approach. It works by selecting a pivot element from the array and partitioning the other elements into two sub-arrays according to whether they are less than or greater than the pivot. The process is then applied recursively to the sub-arrays until the entire array becomes sorted. Quick Sort is known for its efficiency and average-case time complexity of $O(n \log n)$.

Here's an explanation of how Quick Sort works:

1. **Choose a Pivot:** Select an element from the array to act as the pivot. The choice of pivot can significantly affect the algorithm's efficiency. Commonly, the first, last, or middle element of the array is chosen as the pivot, but other strategies, like random selection, can also be used.
2. **Partitioning:** Rearrange the elements in the array so that all elements less than the pivot come before it, and all elements greater than the pivot come after it. After partitioning, the pivot is in its correct sorted position. This step also ensures that the pivot's left side contains elements less than or equal to the pivot, and the right side contains elements greater than or equal to the pivot.
3. **Recursion:** Recursively apply the Quick Sort algorithm to the sub-arrays on both sides of the pivot. This means repeating the partitioning process for the left and right sub-arrays. The recursion continues until each sub-array contains only one element (which is already sorted).
4. **Combine:** Once the recursion reaches the base case (sub-arrays with one element), the sorted sub-arrays are combined to produce the final sorted output.

enjoyalgorithms.com



Quick Sort has an average-case time complexity of $O(n \log n)$ and a worst-case time complexity of $O(n^2)$. The worst-case scenario occurs when the chosen pivot is consistently either the smallest or

largest element, resulting in unbalanced partitions. However, its average-case performance is quite good, and it is often faster than other sorting algorithms like Merge Sort and Heap Sort.

Quick Sort is an in-place sorting algorithm, meaning it requires only a small additional amount of memory for its operations, making it memory-efficient. Additionally, Quick Sort is not a stable sorting algorithm, meaning the relative order of equal elements may change during the sorting process.

Overall, Quick Sort is an efficient and widely used sorting algorithm, particularly for large datasets, due to its average-case time complexity and in-place nature. Proper pivot selection techniques and optimizations can help improve its performance even further.

Analysis of Sorting Algorithms

Sorting algorithms can be analyzed based on their time complexity, space complexity, stability, and whether they are comparison-based or non-comparison-based algorithms. Let's briefly discuss each of these aspects:

Time Complexity:

The time complexity of a sorting algorithm represents how its running time grows with the size of the input data (usually denoted as "n," the number of elements to be sorted). Sorting algorithms can be categorized into different time complexity classes:

- $O(n^2)$ Algorithms: These algorithms, like Bubble Sort, Selection Sort, and Insertion Sort, have a quadratic time complexity. They are less efficient and typically not suitable for large datasets.
- $O(n \log n)$ Algorithms: These algorithms, like Merge Sort, Quick Sort, and Heap Sort, have a linearithmic time complexity. They are more efficient and commonly used for sorting large datasets.
- $O(n)$ Algorithms: Some specialized sorting algorithms, like Counting Sort and Radix Sort, achieve linear time complexity under certain conditions. They can be extremely fast but often require specific data properties.

Space Complexity:

The space complexity of a sorting algorithm represents the additional memory it requires to perform the sorting operation. Algorithms can be categorised as in-place or not in-place:

- In-Place Algorithms: These algorithms sort the elements using a constant amount of extra memory, usually by modifying the input array itself. Examples include Bubble Sort, Selection Sort, and Quick Sort.
- Not In-Place Algorithms: These algorithms require additional memory proportional to the size of the input data to perform the sorting. Examples include Merge Sort and Heap Sort.

Stability:

Stability refers to the property of a sorting algorithm that preserves the relative order of equal elements in the sorted output as they appeared in the input. Stable sorting algorithms maintain the original order of equal elements, while unstable algorithms do not guarantee this property. Stability is relevant when sorting by multiple keys or when maintaining the order of equal elements is necessary.

- Stable Algorithms: Examples of stable sorting algorithms include Merge Sort and Insertion Sort.
- Unstable Algorithms: Examples of unstable sorting algorithms include Quick Sort and Heap Sort.

Comparison-based vs. Non-comparison-based:

Sorting algorithms can be broadly categorised into comparison-based and non-comparison-based algorithms:

- **Comparison-based Algorithms:** These algorithms sort elements by comparing pairs of elements to determine their relative order. Most common sorting algorithms fall into this category, such as Bubble Sort, Selection Sort, Merge Sort, Quick Sort, and Heap Sort.
- **Non-comparison-based Algorithms:** These algorithms do not rely on element comparisons to sort data. Instead, they take advantage of specific properties of the data, making them highly efficient in certain scenarios. Examples include Counting Sort and Radix Sort.

Each sorting algorithm has its strengths and weaknesses, and the choice of the most suitable algorithm depends on the specific requirements of the application, the size of the dataset, and the desired performance characteristics. Understanding the analysis of sorting algorithms can help in selecting the most appropriate algorithm for a given problem and optimising the efficiency of sorting operations.